

Group 4 - ADS 506

November 29, 2022

```
[630]: #Import all required packages:
import warnings
warnings.filterwarnings("ignore")
import os
import itertools
import pandas as pd
import numpy as np
import datetime
import statsmodels.api as sm
from pandas.plotting import autocorrelation_plot
import statsmodels.formula.api as smf
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.seasonal import seasonal_decompose
from sklearn.metrics import mean_squared_error
from dateutil.parser import parse
from statsmodels.tsa.api import ExponentialSmoothing, SimpleExpSmoothing, Holt
import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as plt
from prettytable import PrettyTable
%matplotlib inline
```

```
[631]: df = pd.read_excel('/Users/JohnnyBlaze/Website Data Sets/Online Retail.xlsx',
    ↪ parse_dates=[4])
```

```
[632]: df.head()
```

```
[632]:
```

	InvoiceNo	StockCode	Description	Quantity	\
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	
1	536365	71053	WHITE METAL LANTERN	6	
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	

	InvoiceDate	UnitPrice	CustomerID	Country
0	2010-12-01 08:26:00	2.55	17850.0	United Kingdom
1	2010-12-01 08:26:00	3.39	17850.0	United Kingdom

2	2010-12-01 08:26:00	2.75	17850.0	United Kingdom
3	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
4	2010-12-01 08:26:00	3.39	17850.0	United Kingdom

[633]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   InvoiceNo              541909 non-null object
1   StockCode              541909 non-null object
2   Description            540455 non-null object
3   Quantity               541909 non-null int64
4   InvoiceDate            541909 non-null datetime64[ns]
5   UnitPrice              541909 non-null float64
6   CustomerID             406829 non-null float64
7   Country                541909 non-null object
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB
```

[634]: `df.dropna()`

[634]:

	InvoiceNo	StockCode	Description	Quantity	\
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	
1	536365	71053	WHITE METAL LANTERN	6	
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	
...	
541904	581587	22613	PACK OF 20 SPACEBOY NAPKINS	12	
541905	581587	22899	CHILDREN'S APRON DOLLY GIRL	6	
541906	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4	
541907	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4	
541908	581587	22138	BAKING SET 9 PIECE RETROSPOT	3	

	InvoiceDate	UnitPrice	CustomerID	Country
0	2010-12-01 08:26:00	2.55	17850.0	United Kingdom
1	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
2	2010-12-01 08:26:00	2.75	17850.0	United Kingdom
3	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
4	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
...
541904	2011-12-09 12:50:00	0.85	12680.0	France
541905	2011-12-09 12:50:00	2.10	12680.0	France
541906	2011-12-09 12:50:00	4.15	12680.0	France
541907	2011-12-09 12:50:00	4.15	12680.0	France

```
541908 2011-12-09 12:50:00      4.95      12680.0      France
```

```
[406829 rows x 8 columns]
```

```
[635]: df = df[df['Quantity'] > 0 ]
```

```
[636]: len(df)
```

```
[636]: 531285
```

```
[637]: df['Sales'] = (df['Quantity'] * df['UnitPrice'])
```

```
[638]: df.head()
```

```
[638]: InvoiceNo StockCode      Description  Quantity \
0    536365    85123A  WHITE HANGING HEART T-LIGHT HOLDER      6
1    536365    71053           WHITE METAL LANTERN      6
2    536365    84406B    CREAM CUPID HEARTS COAT HANGER      8
3    536365    84029G  KNITTED UNION FLAG HOT WATER BOTTLE      6
4    536365    84029E    RED WOOLLY HOTTIE WHITE HEART.      6
```

```
InvoiceDate  UnitPrice  CustomerID      Country  Sales
0 2010-12-01 08:26:00      2.55    17850.0  United Kingdom  15.30
1 2010-12-01 08:26:00      3.39    17850.0  United Kingdom  20.34
2 2010-12-01 08:26:00      2.75    17850.0  United Kingdom  22.00
3 2010-12-01 08:26:00      3.39    17850.0  United Kingdom  20.34
4 2010-12-01 08:26:00      3.39    17850.0  United Kingdom  20.34
```

```
[639]: clock = df[df['Description'].str.contains('CLOCK', na=False)]
```

```
[640]: len(clock)
```

```
[640]: 7180
```

```
[641]: clock['Category'] = 'Clock'
```

```
[642]: clock.drop(['InvoiceNo', 'StockCode', 'Description', 'CustomerID', 'Quantity',
↪ 'UnitPrice'], axis=1, inplace=True)
```

```
[643]: clock.head()
```

```
[643]: InvoiceDate      Country  Sales  Category
26 2010-12-01 08:45:00      France   90.0    Clock
27 2010-12-01 08:45:00      France   90.0    Clock
28 2010-12-01 08:45:00      France   45.0    Clock
149 2010-12-01 09:45:00  United Kingdom   15.0    Clock
204 2010-12-01 10:03:00      Australia   17.0    Clock
```

```
[644]: clock['Country'].value_counts()
```

```
[644]: United Kingdom      6415
      France              184
      Germany             141
      EIRE                127
      Belgium             63
      Australia           38
      Switzerland         35
      Denmark             25
      Spain               24
      Norway              21
      Iceland             18
      Channel Islands     13
      Netherlands         13
      Portugal            11
      Finland             11
      Singapore           7
      Cyprus              7
      Israel              6
      Canada              5
      Poland              3
      Greece              2
      Brazil              2
      Malta               2
      Italy               2
      RSA                 2
      Hong Kong           1
      European Community  1
      Unspecified         1
      Name: Country, dtype: int64
```

```
[645]: clock = clock.loc[clock['Country'] == 'United Kingdom']
```

```
[646]: clock.drop('Country', axis=1, inplace=True)
```

```
[647]: len(clock)
```

```
[647]: 6415
```

```
[648]: clock.head()
```

```
[648]:      InvoiceDate  Sales  Category
149  2010-12-01 09:45:00   15.0    Clock
271  2010-12-01 10:47:00   15.0    Clock
272  2010-12-01 10:47:00   30.0    Clock
273  2010-12-01 10:47:00   30.0    Clock
274  2010-12-01 10:47:00   30.0    Clock
```

```
[649]: type(clock)
```

```
[649]: pandas.core.frame.DataFrame
```

```
[650]: # clock3 = pd.DataFrame(clock.groupby("InvoiceDate")['Sales'].sum().  
      ↪reset_index())  
      # clock3.head()
```

```
[651]: # clock3.plot(figsize=(12, 4))  
      # plt.title('Clock Sales', fontweight='bold', size=20)  
      # plt.show()
```

```
[652]: # clock.resample('H', on='InvoiceDate').Sales.sum()
```

```
[653]: clock2 = clock.resample('D', on='InvoiceDate').Sales.sum().reset_index()
```

```
[654]: clock2 = pd.DataFrame(clock2)
```

```
[655]: clock2.head()
```

```
[655]: InvoiceDate  Sales  
0  2010-12-01  568.40  
1  2010-12-02  798.25  
2  2010-12-03  587.62  
3  2010-12-04    0.00  
4  2010-12-05  596.00
```

```
[656]: clock2 = clock2[clock2['Sales'] > 1]
```

```
[657]: type(clock2)
```

```
[657]: pandas.core.frame.DataFrame
```

```
[658]: len(clock2)
```

```
[658]: 303
```

```
[659]: clock2 = clock2.rename(columns={'Sales': 'Clock_Sales'})
```

```
[660]: clock2.head()
```

```
[660]: InvoiceDate  Clock_Sales  
0  2010-12-01    568.40  
1  2010-12-02    798.25  
2  2010-12-03    587.62  
4  2010-12-05    596.00  
5  2010-12-06    475.97
```

```
[661]: clock2.tail()
```

```
[661]: InvoiceDate  Clock_Sales
369  2011-12-05      1246.29
370  2011-12-06      1358.44
371  2011-12-07      2496.14
372  2011-12-08      1195.32
373  2011-12-09       421.82
```

```
[662]: clock2['InvoiceDate'].min()
```


```
[662]: Timestamp('2010-12-01 00:00:00')
```

```
[663]: clock2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 303 entries, 0 to 373
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   InvoiceDate      303 non-null   datetime64[ns]
1   Clock_Sales     303 non-null   float64
dtypes: datetime64[ns](1), float64(1)
memory usage: 7.1 KB
```


```
[664]: def plot_df(clock2, x, y, title="", xlabel='Date', ylabel='# of Clocks Sold',
↳ dpi=100):
    plt.figure(figsize=(15,4), dpi=dpi)
    plt.plot(x, y, color='tab:green')
    plt.gca().set(title=title, xlabel=xlabel, ylabel=ylabel)
    plt.legend(['Clock_Sales'])
    plt.show()

plot_df(df, x=clock2['InvoiceDate'], y=clock2['Clock_Sales'], title='Online_
↳ Clock Sales Dec 2010 to Dec 2011')
```



```
[665]: x = clock2['InvoiceDate'].values
y1 = clock2['Clock_Sales'].values

# Plot
fig, ax = plt.subplots(1, 1, figsize=(16,5), dpi= 120)
plt.fill_between(x, y1=y1, y2=-y1, alpha=0.5, linewidth=2, color='blue')
plt.ylim(-6000, 6000)
plt.title('Clock Sales (Two Side View)', fontsize=16)
plt.hlines(y=0, xmin=np.min(clock2['InvoiceDate']), xmax=np.
    ↪max(clock2['InvoiceDate']), linewidth=.5)
plt.show()
```



```
[666]: # Decomposition
# Decomposition of a time series can be performed by considering the series as
    ↪an additive or multiplicative combination of the base level, trend, seasonal
    ↪index and the residual term.

from statsmodels.tsa.seasonal import seasonal_decompose
from dateutil.parser import parse


# Multiplicative Decomposition
multiplicative_decomposition = seasonal_decompose(clock2['Clock_Sales'],
    ↪model='multiplicative', period=35)


# Additive Decomposition
additive_decomposition = seasonal_decompose(clock2['Clock_Sales'],
    ↪model='additive', period=35)

# Plot
plt.rcParams.update({'figure.figsize': (16,12)})
multiplicative_decomposition.plot().suptitle('Multiplicative Decomposition',
    ↪fontsize=16)
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
```

```
additive_decomposition.plot().subtitle('Additive Decomposition', fontsize=16)  
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
```

```
plt.show()
```

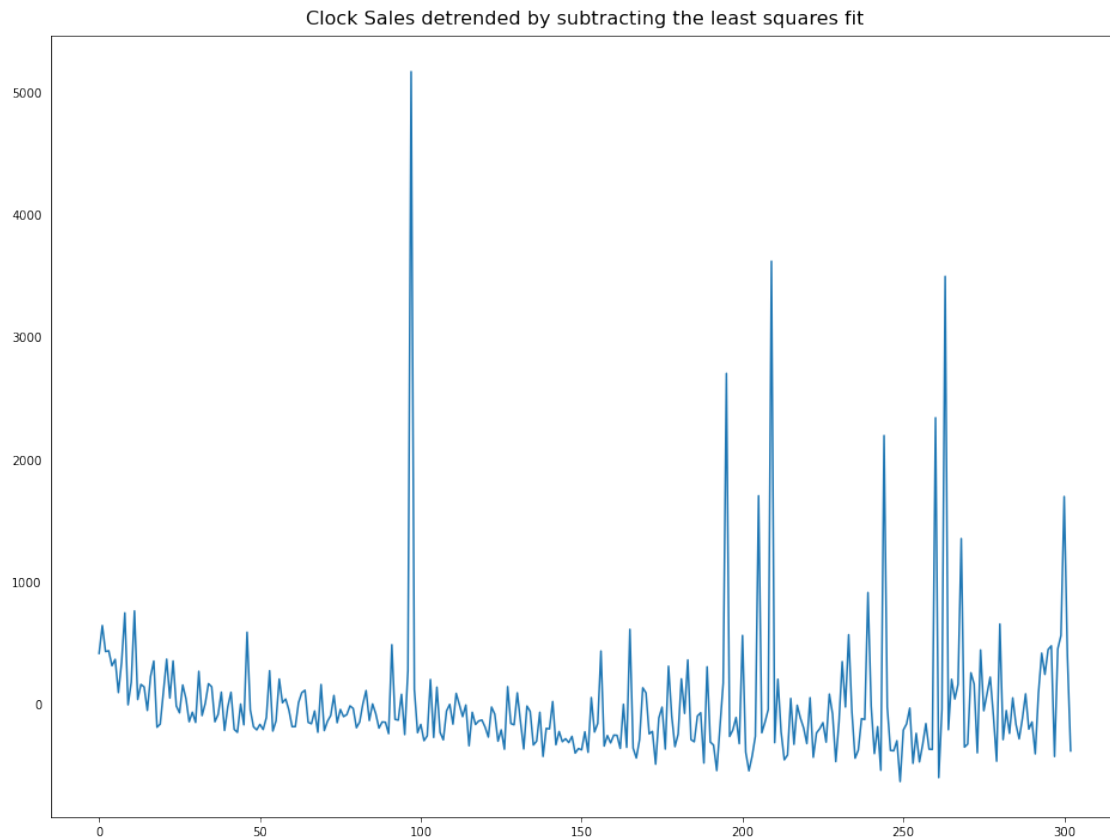




- 0.1 If we look at the residuals of the additive decomposition closely, it has some pattern left over.
- 0.2 The multiplicative decomposition, looks quite random which is good. So ideally, multiplicative decomposition should be preferred for this particular series.

```
[667]: # Detrend

from scipy import signal
detrended = signal.detrend(clock2['Clock_Sales'].values)
plt.plot(detrended)
plt.title('Clock Sales detrended by subtracting the least squares fit',
         ↪fontsize=16)
plt.show()
```



0.3 The graph looks the same as when it was first plotted suggesting no trend

```
[668]: # Dickey-Fuller Test
from statsmodels.tsa.stattools import adfuller

#perform augmented Dickey-Fuller test
X = clock2['Clock_Sales'].values
result = adfuller(X)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
print('Critical Values:')
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))
```

```
ADF Statistic: -4.091994
p-value: 0.000998
Critical Values:
1%: -3.453
5%: -2.871
10%: -2.572
```

- 0.4 Running the example prints the test statistic value of -4. The more negative this statistic, the more likely we are to reject the null hypothesis (we have a stationary dataset).
- 0.5 This suggests that we can reject the null hypothesis with a significance level of less than 1% (i.e. a low probability that the result is a statistical fluke).
- 0.6 Rejecting the null hypothesis means that the process has no unit root, and in turn that the time series is stationary or does not have time-dependent structure.
- 0.7 Stationarity of the time-series data: The stationarity of the data can be found using `adfuller` class of `statsmodels.tsa.stattools` module. The value of p-value is used to determine whether there is stationarity. If the value is less than 0.05, the stationarity exists.

```
[669]: # Dickey-Fuller Test

from statsmodels.tsa.stattools import adfuller

# Run the test


df_stationarityTest = adfuller(clock2['Clock_Sales'], autolag='AIC')

# Check the value of p-value

print("P-value: ", df_stationarityTest[1])
```


P-value: 0.000998459148273761

```
[670]: clock2['Clock_Sales'].hist()
plt.show()
```



```
[671]: # Test for seasonality
from pandas.plotting import autocorrelation_plot

# Draw Plot
plt.rcParams.update({'figure.figsize':(10,8), 'figure.dpi':100})
autocorrelation_plot(clock2['Clock_Sales'].tolist())
plt.show()
```




- 0.8 If partial autocorrelation values are close to 0, then values between observations and lagged observations are not correlated with one another. Inversely, partial autocorrelations with values close to 1 or -1 indicate that there exists strong positive or negative correlations between the lagged observations of the time series.

```
[672]: # Lag Plots

from pandas.plotting import lag_plot
plt.rcParams.update({'ytick.left' : False, 'axes.titlepad':10})

# Plot
fig, axes = plt.subplots(1, 4, figsize=(10,3), sharex=True, sharey=True,
    dpi=100)
for i, ax in enumerate(axes.flatten()[:4]):
    lag_plot(clock2['Clock_Sales'], lag=i+1, ax=ax, c='seagreen')
    ax.set_title('Lag ' + str(i+1))


fig.suptitle('Lag Plots of Clock Sales', y=1.05)
plt.show()
```



0.9 A Lag plot is a scatter plot of a time series against a lag of itself. It is normally used to check for autocorrelation. If there is any pattern existing in the series, the series is autocorrelated. If there is no such pattern, the series is likely to be random white noise.

```
[673]: from statsmodels.graphics.tsaplots import plot_pacf

pacf = plot_pacf(clock2['Clock_Sales'], lags=25)
```



0.10 The above plot can be used to determine the order of AR model. You may note that a correlation value up to order 3 is high enough. Thus, we will train the AR model of order 3.

0.10.1 Naive Forecast Model

```
[674]: # Split Train / Test

train_length = 243
train = clock2[0:train_length]
test = clock2[train_length:]
print(len(train))
print('')
print(len(test))
```


243

60

```
[675]: # Naive Forecast
```

```
naive = test.copy()
naive['naive_forecast'] = train['Clock_Sales'][train_length-1]

plt.figure(figsize=(20,5))
plt.grid()
plt.plot(train['Clock_Sales'], label='Train')
plt.plot(test['Clock_Sales'], label='Test')
plt.plot(naive['naive_forecast'], label='Naive forecast')
plt.legend(loc='best')
plt.title('Naive Method')
plt.show()
```



```
[676]: n_rmse = np.sqrt(mean_squared_error(test['Clock_Sales'],
↳naive['naive_forecast'])).round(2)
n_mape = np.round(np.mean(np.abs(test['Clock_Sales']-naive['naive_forecast'])/
↳test['Clock_Sales'])*100,2)

results = pd.DataFrame({'Method':['Naive method'], 'MAPE': [mape], 'RMSE':
↳[rmse]})
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
[676]:
```

	Method	RMSE	MAPE
0	Naive method	868.59	58.44

Per the graph naive method is not suitable for data with high variability


0.10.2 Simple Average

```
[677]: simple_average = test.copy()
simple_average['avg_forecast'] = train['Clock_Sales'].mean()

plt.figure(figsize=(20,5))
```



```
plt.grid()
plt.plot(train['Clock_Sales'], label='Train')
plt.plot(test['Clock_Sales'], label='Test')
plt.plot(simple_average['avg_forecast'], label='Average forecast')
plt.legend(loc='best')
plt.title('Average Method')
plt.show()
```



```
[678]: sa_rmse = np.sqrt(mean_squared_error(test['Clock_Sales'],
↪simple_average['avg_forecast'])).round(2)
sa_mape = np.round(np.mean(np.
↪abs(test['Clock_Sales']-simple_average['avg_forecast'])/
↪test['Clock_Sales'])*100,2)

results = pd.DataFrame({'Method':['Average method'], 'MAPE': [mape], 'RMSE':
↪[rmse]})
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
[678]:           Method    RMSE    MAPE
0  Average method  868.59  58.44
```


This model did improve our score, it seems the average of our data is pretty consistent.

0.10.3 Moving Average (60 Day)

```
[679]: moving_avg = test.copy()
moving_avg['moving_avg_forecast'] = train['Clock_Sales'].rolling(60).mean().
↪iloc[-1]

plt.figure(figsize=(20,5))
plt.grid()
plt.plot(train['Clock_Sales'], label='Train')
plt.plot(test['Clock_Sales'], label='Test')
plt.plot(moving_avg['moving_avg_forecast'], label='Moving Average Forecast')
```

```
plt.legend(loc='best')
plt.show()
```



```
[680]: ma_rmse = np.sqrt(mean_squared_error(test['Clock_Sales'],
      ↪moving_avg['moving_avg_forecast'])).round(2)
ma_mape = np.round(np.mean(np.
      ↪abs(test['Clock_Sales']-moving_avg['moving_avg_forecast'])/
      ↪test['Clock_Sales'])*100,2)

results = pd.DataFrame({'Method':['Moving Average method'], 'MAPE': [mape],
      ↪'RMSE': [rmse]})
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
[680]:
```


	Method	RMSE	MAPE
0	Moving Average method	868.59	58.44

Interestingly enough this model did not improve our results after choosing the last 60 days. We could adjust the window and see if that improves our results.

0.10.4 Simple Exponential Smoothing

```
[681]: ses = test.copy()
ses_fit = SimpleExpSmoothing(np.asarray(train['Clock_Sales'])).
      ↪fit(smoothing_level=0.6,optimized=False)
ses['SES'] = ses_fit.forecast(len(test))

plt.figure(figsize=(20,5))
plt.grid()
plt.plot(train['Clock_Sales'], label='Train')
plt.plot(test['Clock_Sales'], label='Test')
plt.plot(ses['SES'], label='SES')
plt.legend(loc='best')
plt.show()
```



```
[682]: se_rmse = np.sqrt(mean_squared_error(test['Clock_Sales'], ses['SES'])).round(2)
se_mape = np.round(np.mean(np.abs(test['Clock_Sales']-ses['SES'])/
    ↪test['Clock_Sales'])*100,2)

results = pd.DataFrame({'Method':['Simple Exponential Smoothing method'],
    ↪'MAPE': [mape], 'RMSE': [rmse]})
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
[682]:
```


	Method	RMSE	MAPE
0	Simple Exponential Smoothing method	868.59	58.44

So far the second best model after simple average. We can tune to alpha from 0.6 to another number to see if it helps improve the model.

0.10.5 Holt Method

```
[683]: holt = test.copy()
holt_fit = Holt(np.asarray(train['Clock_Sales'])).fit(smoothing_level = 0.3,
    ↪smoothing_slope = 0.1)
holt['Holt_linear'] = holt_fit.forecast(len(test))

plt.figure(figsize=(20,5))
plt.grid()
plt.plot(train['Clock_Sales'], label='Train')
plt.plot(test['Clock_Sales'], label='Test')
plt.plot(holt['Holt_linear'], label='Holt_linear')
plt.legend(loc='best')
plt.show()
```



```
[684]: hl_rmse = np.sqrt(mean_squared_error(test['Clock_Sales'], holt['Holt_linear'])).
        ↪round(2)
hl_mape = np.round(np.mean(np.abs(test['Clock_Sales']-holt['Holt_linear'])/
        ↪test['Clock_Sales'])*100,2)

results = pd.DataFrame({'Method':['Holt Linear method'], 'MAPE': [mape], 'RMSE':
        ↪ [rmse]})
results = results[['Method', 'RMSE', 'MAPE']]
results
```


```
[684]:           Method    RMSE    MAPE
0  Holt Linear method  868.59  58.44
```

Results were not very good on the first run, model can be tuned to see if there's improvement

0.10.6 Holt Winter's Method

```
[685]: hw = test.copy()
hw_fit = ExponentialSmoothing(np.asarray(train['Clock_Sales']))
        ↪,seasonal_periods=7 ,trend='add', seasonal='add',).fit()
hw['Holt_Winter'] = hw_fit.forecast(len(test))

plt.figure(figsize=(20,5))
plt.grid()
plt.plot( train['Clock_Sales'], label='Train')
plt.plot(test['Clock_Sales'], label='Test')
plt.plot(hw['Holt_Winter'], label='Holt_Winter')
plt.legend(loc='best')
plt.show()
```



```
[686]: hw_rmse = np.sqrt(mean_squared_error(test['Clock_Sales'], hw['Holt_Winter'])).
        ↪round(2)
hw_mape = np.round(np.mean(np.abs(test['Clock_Sales']-hw['Holt_Winter'])/
        ↪test['Clock_Sales'])*100,2)


results = pd.DataFrame({'Method':['Holt Winters method'], 'MAPE': [mape],
        ↪'RMSE': [rmse]})
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
[686]:
```


	Method	RMSE	MAPE
0	Holt Winters method	868.59	58.44

0.10.7 Arima Model

```
[687]: autocorrelation_plot(clock2['Clock_Sales'])
plt.show()
```



```
[688]: fig = plt.figure(figsize=(12,8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(clock2['Clock_Sales'],lags=40,ax=ax1)
ax2 = fig.add_subplot(212)
fig = sm.graphics.tsa.plot_pacf(clock2['Clock_Sales'],lags=40,ax=ax2)
```



```
[689]: from statsmodels.tsa.arima_model import ARIMA

model = ARIMA(clock2['Clock_Sales'], order = (1,0,0))
model_fit = model.fit(dispatch=0)
model_fit.summary()
```

```
[689]: <class 'statsmodels.iolib.summary.Summary'>
      """
              ARMA Model Results
=====
Dep. Variable:          Clock_Sales    No. Observations:          303
Model:                  ARMA(1, 0)    Log Likelihood             -2368.473
Method:                 css-mle       S.D. of innovations        600.512
Date:                   Tue, 29 Nov 2022    AIC                       4742.946
Time:                   21:22:27          BIC                       4754.087
Sample:                 0              HQIC                      4747.403

=====
=====
      coef    std err          z      P>|z|      [0.025
0.975]
-----
-----
```

```

const          481.2924    38.920    12.366    0.000    405.011
557.574
ar.L1.Clock_Sales    0.1140    0.057    2.000    0.045    0.002
0.226

```


Roots				
	Real	Imaginary	Modulus	Frequency
AR.1	8.7734	+0.0000j	8.7734	0.0000

```

"""

```

```
[690]: model_fit.plot_predict(dynamic=False)
plt.show()
```



```
[691]: # Create Training and Test For Arima
train_length = 243

train_sales = clock2.Clock_Sales[0:train_length]
test_sales = clock2.Clock_Sales[train_length:]
```




```
[692]: # Build Model
model = ARIMA(train_sales, order=(1, 0, 0))
fitted = model.fit(dispatch=0)

# Forecast using 95% confidence interval
fc, se, conf = fitted.forecast(60, alpha=0.05)

# Make as pandas series
fc_series = pd.Series(fc, index=test_sales.index)
lower_series = pd.Series(conf[:, 0], index=test_sales.index)
upper_series = pd.Series(conf[:, 1], index=test_sales.index)

# Plot
plt.figure(figsize=(12,5), dpi=100)
plt.plot(train_sales, label='training')
plt.plot(test_sales, label='actual')
plt.plot(fc_series, label='forecast')
plt.fill_between(lower_series.index, lower_series, upper_series, color='k',
    ↪alpha=.15)
plt.title('Forecast vs Actuals')
plt.legend(loc='upper left', fontsize=8)
plt.show()
```



```
[693]: # Calculating RMSE and MAPE

arima_rmse = np.sqrt(mean_squared_error(test_sales, fc_series)).round(2)
arima_mape = np.round(np.mean(np.abs(test_sales - fc_series)/test_sales)*100,2)

results = pd.DataFrame({'Method':['ARIMA method'], 'MAPE': [mape], 'RMSE':
    ↪[rmse]})
```

```
results = results[['Method', 'RMSE', 'MAPE']]
results
```

```
[693]:      Method      RMSE   MAPE
0  ARIMA method  868.59  58.44
```

```
[695]: # Table Results
```

```
Table = PrettyTable(["Model", "MAPE", "RMSE"])
Table.add_row(["Naive", n_mape, n_rmse])
Table.add_row(["Simple Average", sa_mape, sa_rmse])
Table.add_row(["Moving Average", ma_mape, ma_rmse])
Table.add_row(["Simple Exponential", se_mape, se_rmse])
Table.add_row(["Holt Linear", hl_mape, hl_rmse])
Table.add_row(["Holt Winter", hw_mape, hw_rmse])
Table.add_row(["ARIMA", arima_mape, arima_rmse])
print("Time Series Model Performance Sorted by MAPE")
Table.sortby = "MAPE"
print(Table)
```

Time Series Model Performance Sorted by MAPE

Model	MAPE	RMSE
Simple Average	58.35	868.61
ARIMA	58.44	868.59
Simple Exponential	63.33	819.63
Naive	65.53	962.45
Moving Average	74.06	778.11
Holt Linear	74.44	776.13
Holt Winter	83.05	729.28